The Meaning of Uniform Manifold Approximation and Projection (UMAP)

Nick Lines

JOHNS HOPKINS UNIVERSITY

https://github.com/linesn/the_meaning_of_umap

What is UMAP?

- Non-linear dimensionality reduction
- Invented in 2018 by
  Leland McInnes
  John Healy
  James Melville
  (Tutte Institute) [1]

UNIFORM MANIFOLD
UMAP
APPROXIMATION & PROJECTION

Outline

1. What does UMAP do?
2. How does it do that?
3. Practical considerations
4. Conclusion

Remember: Nothing I'll say here is original!

1. What does UMAP do?

Guernica, by Picasso

Guernica, by Picasso

"I make the painting for the painting. I paint the objects for what they are." [4]

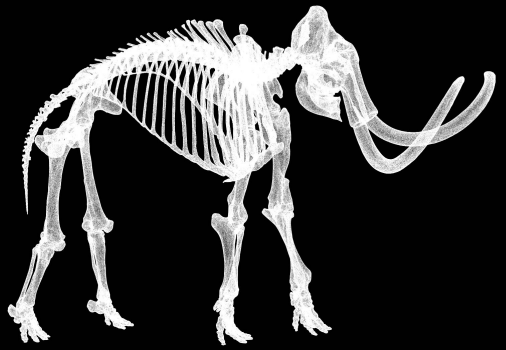

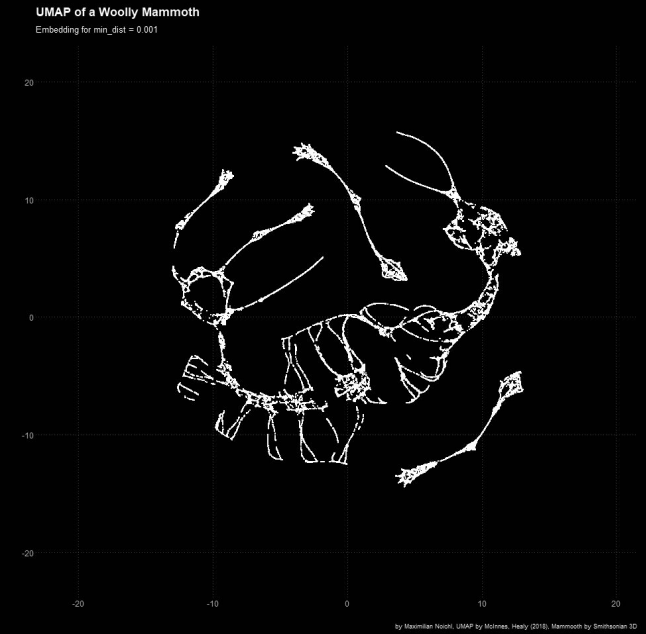

Examples of UMAP Dimensionality Reduction applied to Smithsonian 3D models [6]

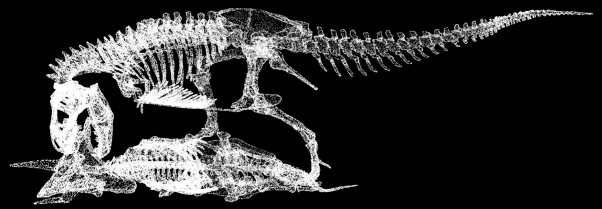

2. How does UMAP do that?

Dimensionality Reduction

Matrix Factorization

Neighborhood Graphs

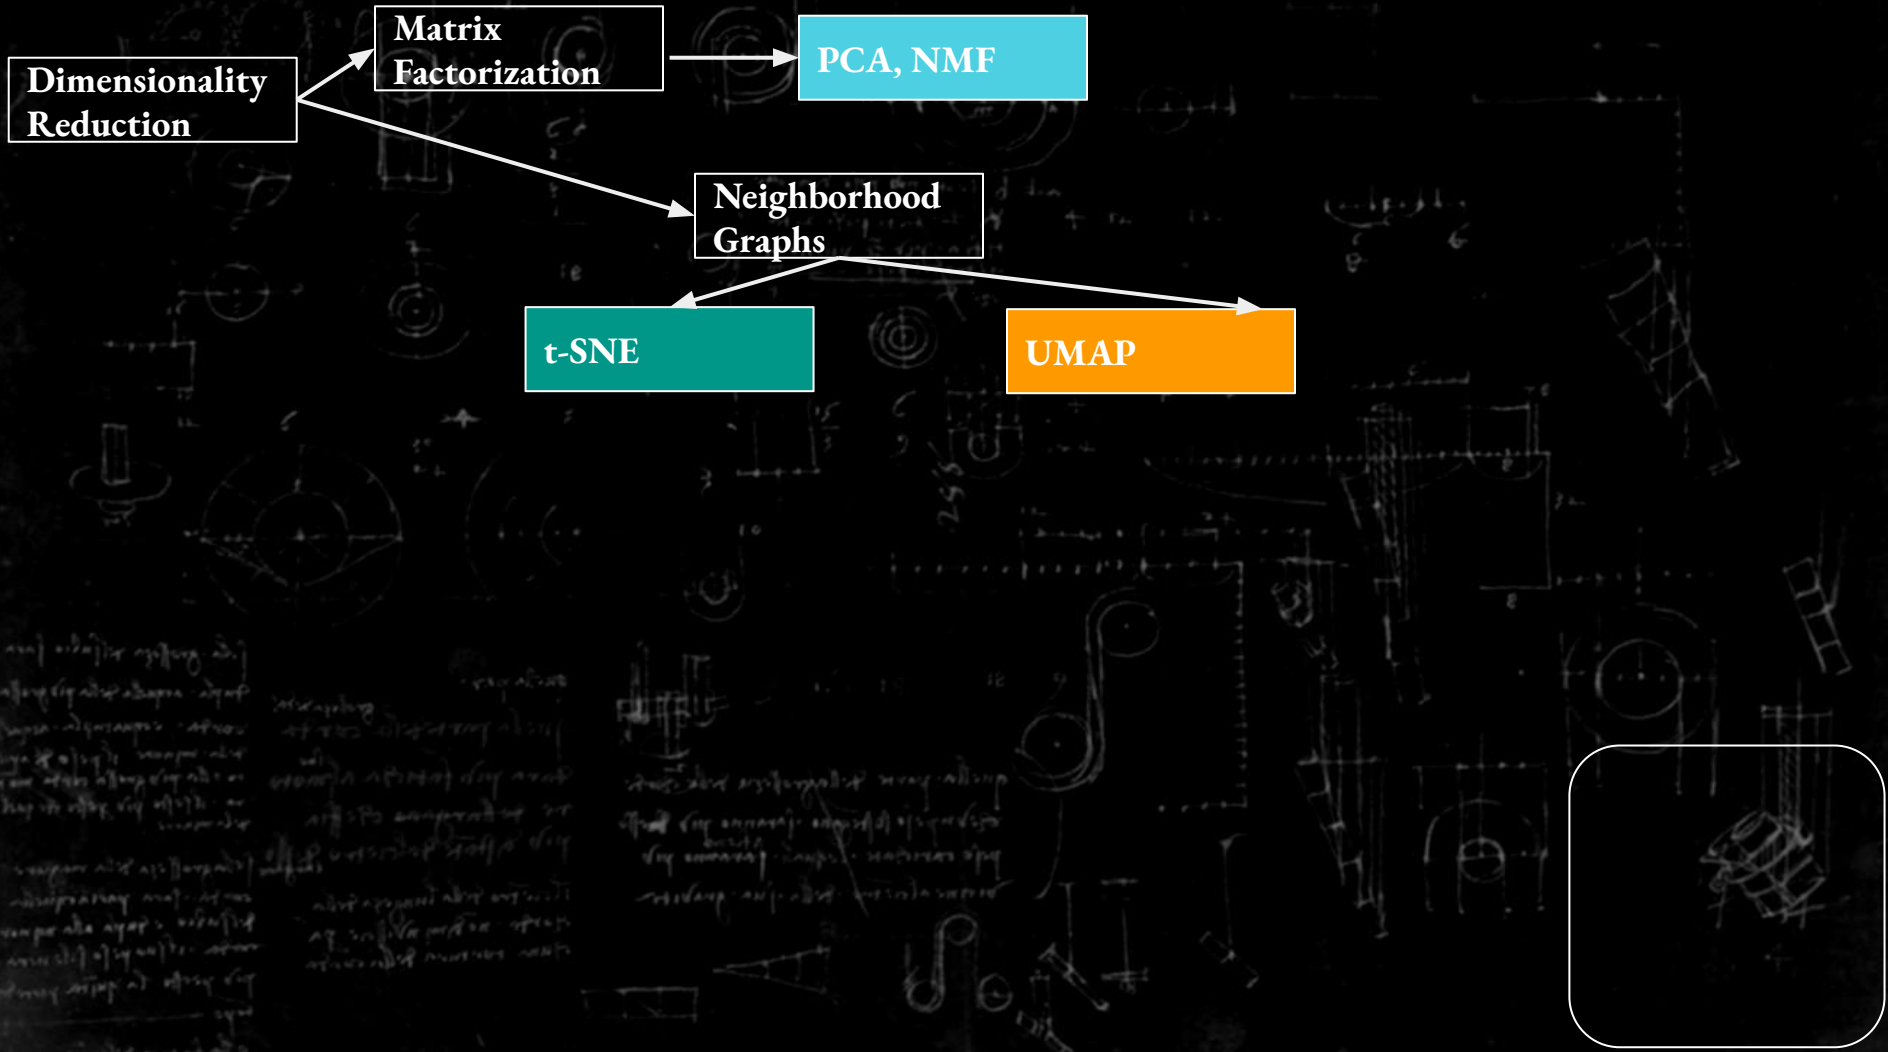
Dimensionality Reduction
Matrix Factorization
PCA, NMF
Neighborhood Graphs
t-SNE
UMAP

Dimensionality Reduction

Matrix Factorization

Neighborhood Graphs

t-SNE

UMAP

Pairwise Distances

Convert to Probability Dists.

Minimize Kullback-Leibler with Gradient Descent

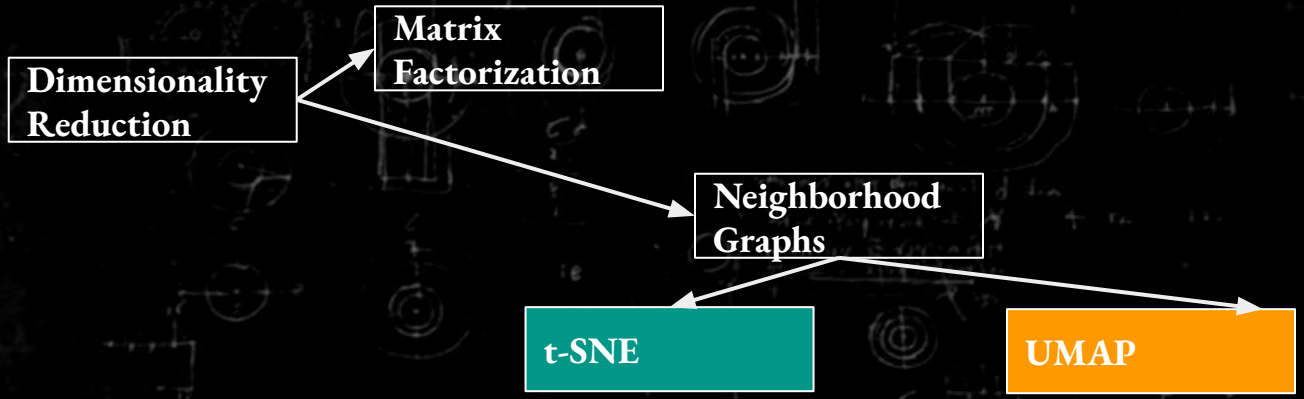

Find the topological structure in high dimension representation

Pairwise Distances

Convert to Probability Dists.

Initialize then optimize a lower dimensional representation

Minimize Kullback-Leibler with Gradient Descent

Dimensionality Reduction

Matrix Factorization

Neighborhood Graphs

t-SNE

UMAP

Pairwise Distances

Perplexity

Convert to Probability Dists.

Minimize Kullback-Leibler with Gradient Descent

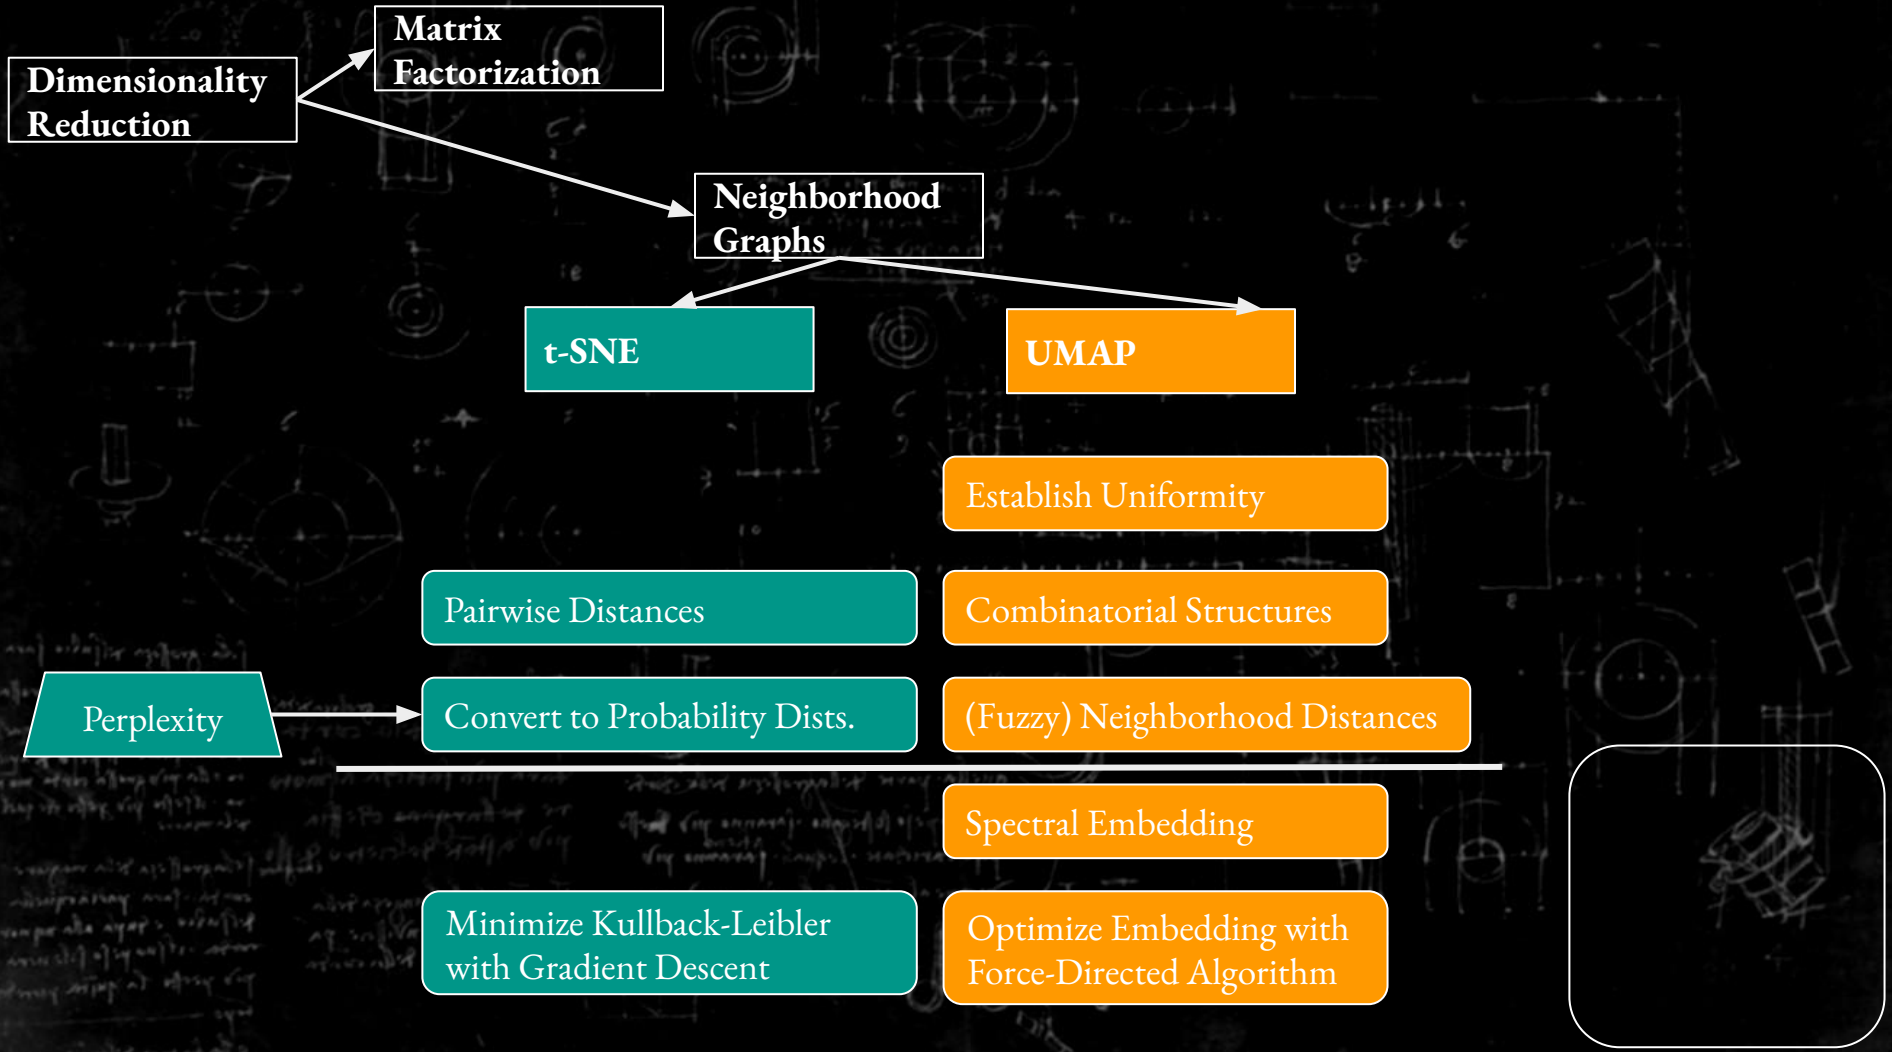
Dimensionality Reduction
Matrix Factorization
Neighborhood Graphs
t-SNE
UMAP
Establish Uniformity
Pairwise Distances
Combinatorial Structures
Perplexity
Convert to Probability Dists.
(Fuzzy) Neighborhood Distances
Spectral Embedding
Minimize Kullback-Leibler with Gradient Descent
Optimize Embedding with Force-Directed Algorithm

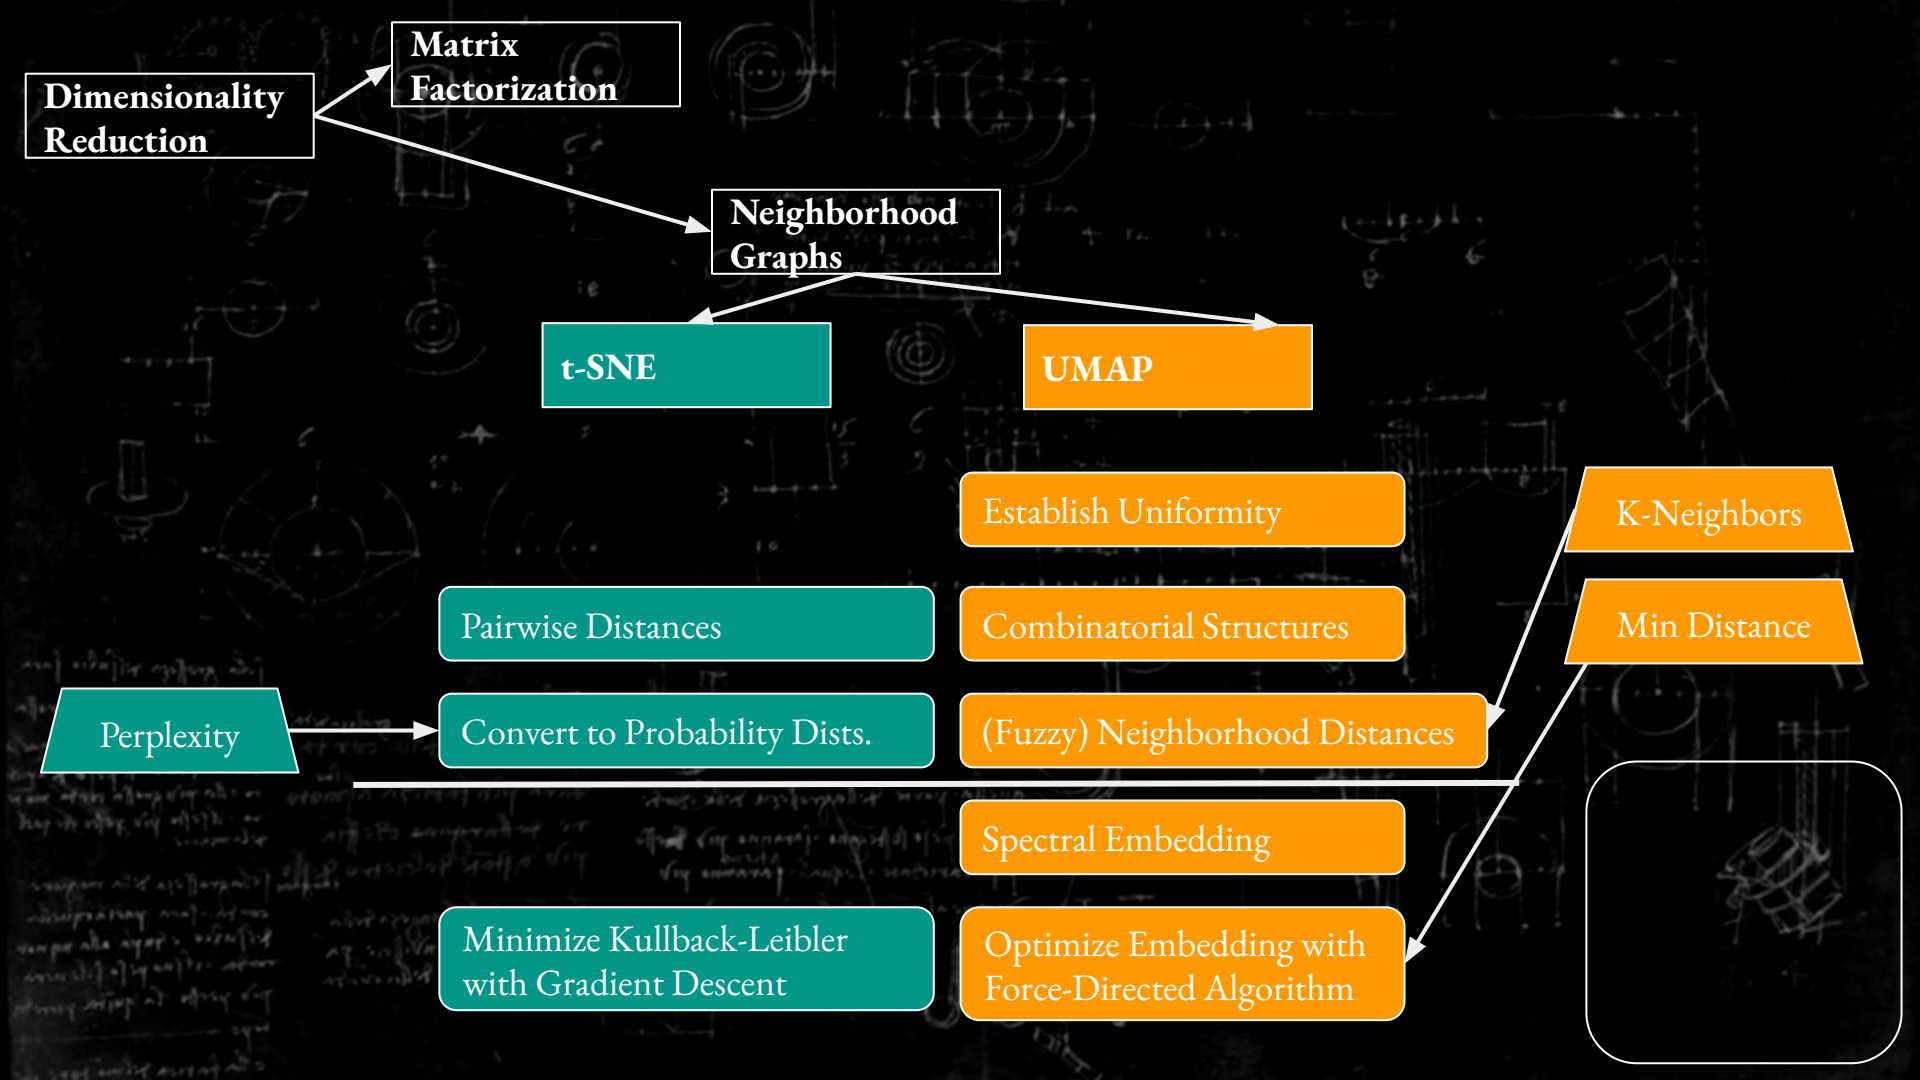
Dimensionality Reduction
Matrix Factorization
Neighborhood Graphs
t-SNE
UMAP
Establish Uniformity
K-Neighbors
Pairwise Distances
Combinatorial Structures
Min Distance
Perplexity
Convert to Probability Dists.
(Fuzzy) Neighborhood Distances
Spectral Embedding
Minimize Kullback-Leibler with Gradient Descent
Optimize Embedding with Force-Directed Algorithm

Justify the assumption “uniformly distributed” data

(Builds on David Spivak’s work [5])

$$\boldsymbol{x} = x_1, x_2, \ldots, x_r$$

- To form consistent local neighborhoods that define graph structure, we want uniformly distributed data.
- NOT true by default with finite samples in a high dimension manifold
- So establish a distance function d_i for each x_i so that the unit ball centered at x_i contains the nearest k data points.

Switch to convenient combinatorial structures

[2]

- A *k-simplex* is a structure formed by *k+1* (independent) points in *k* dimensions.
- A *simplicial complex* is a set *K* of simplices such that each face of each simplex in *K* is in *K*, and the (nonempty) intersection of any two simplicies in *K* is a shared face.

- So we build a simplicial complex called the Čech complex by describing each data point as a 0-simplex. Then we connect each point to other points within the unit ball around it, and build 1-simplices, then 2-simplices, etc. until all points are part of the complex.

- The **Nerve Theorem** guarantees that this preserves the topological structure of the data!

Fuzzy neighborhood distances

- Probability of an edge x_i to x_j in the sense of d_i is a function of the distance $d_i(x_i, x_j)$. More distant points are less likely to be connected.
- Probability that x_i and x_j share an edge is the probability at least one of (x_i to x_j and x_j to x_i) exists.

Fuzzy neighborhood distances

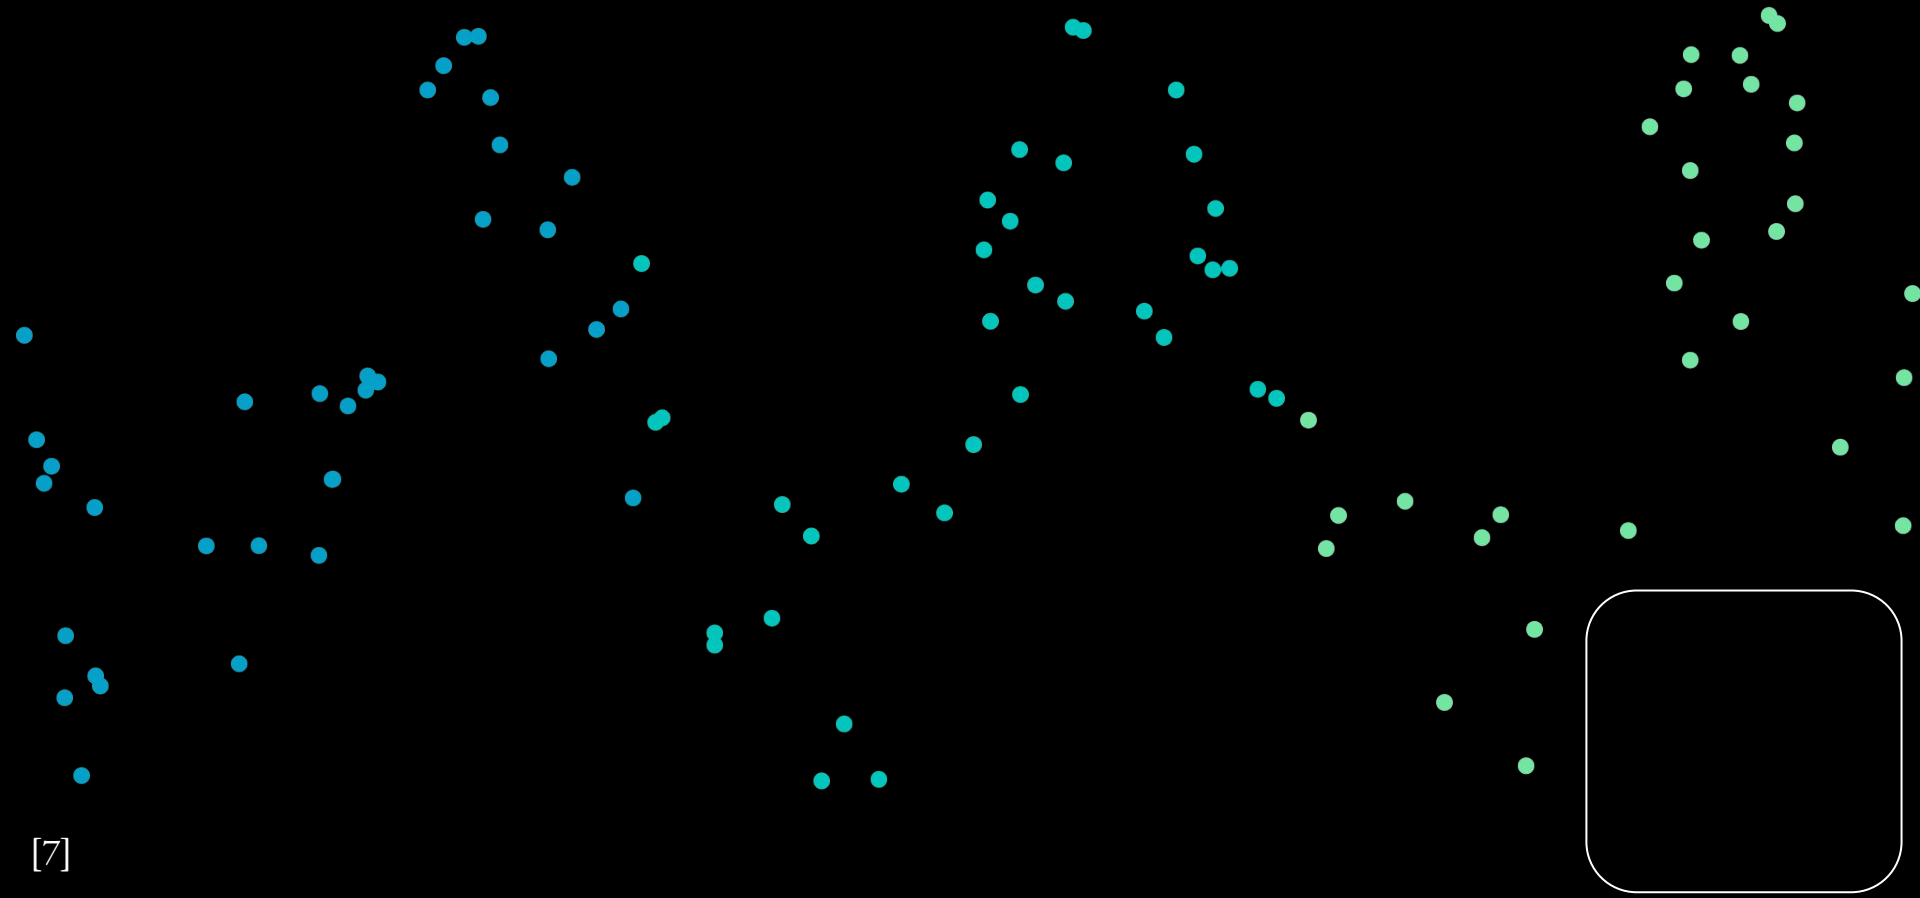

[7]

Fuzzy neighborhood distances

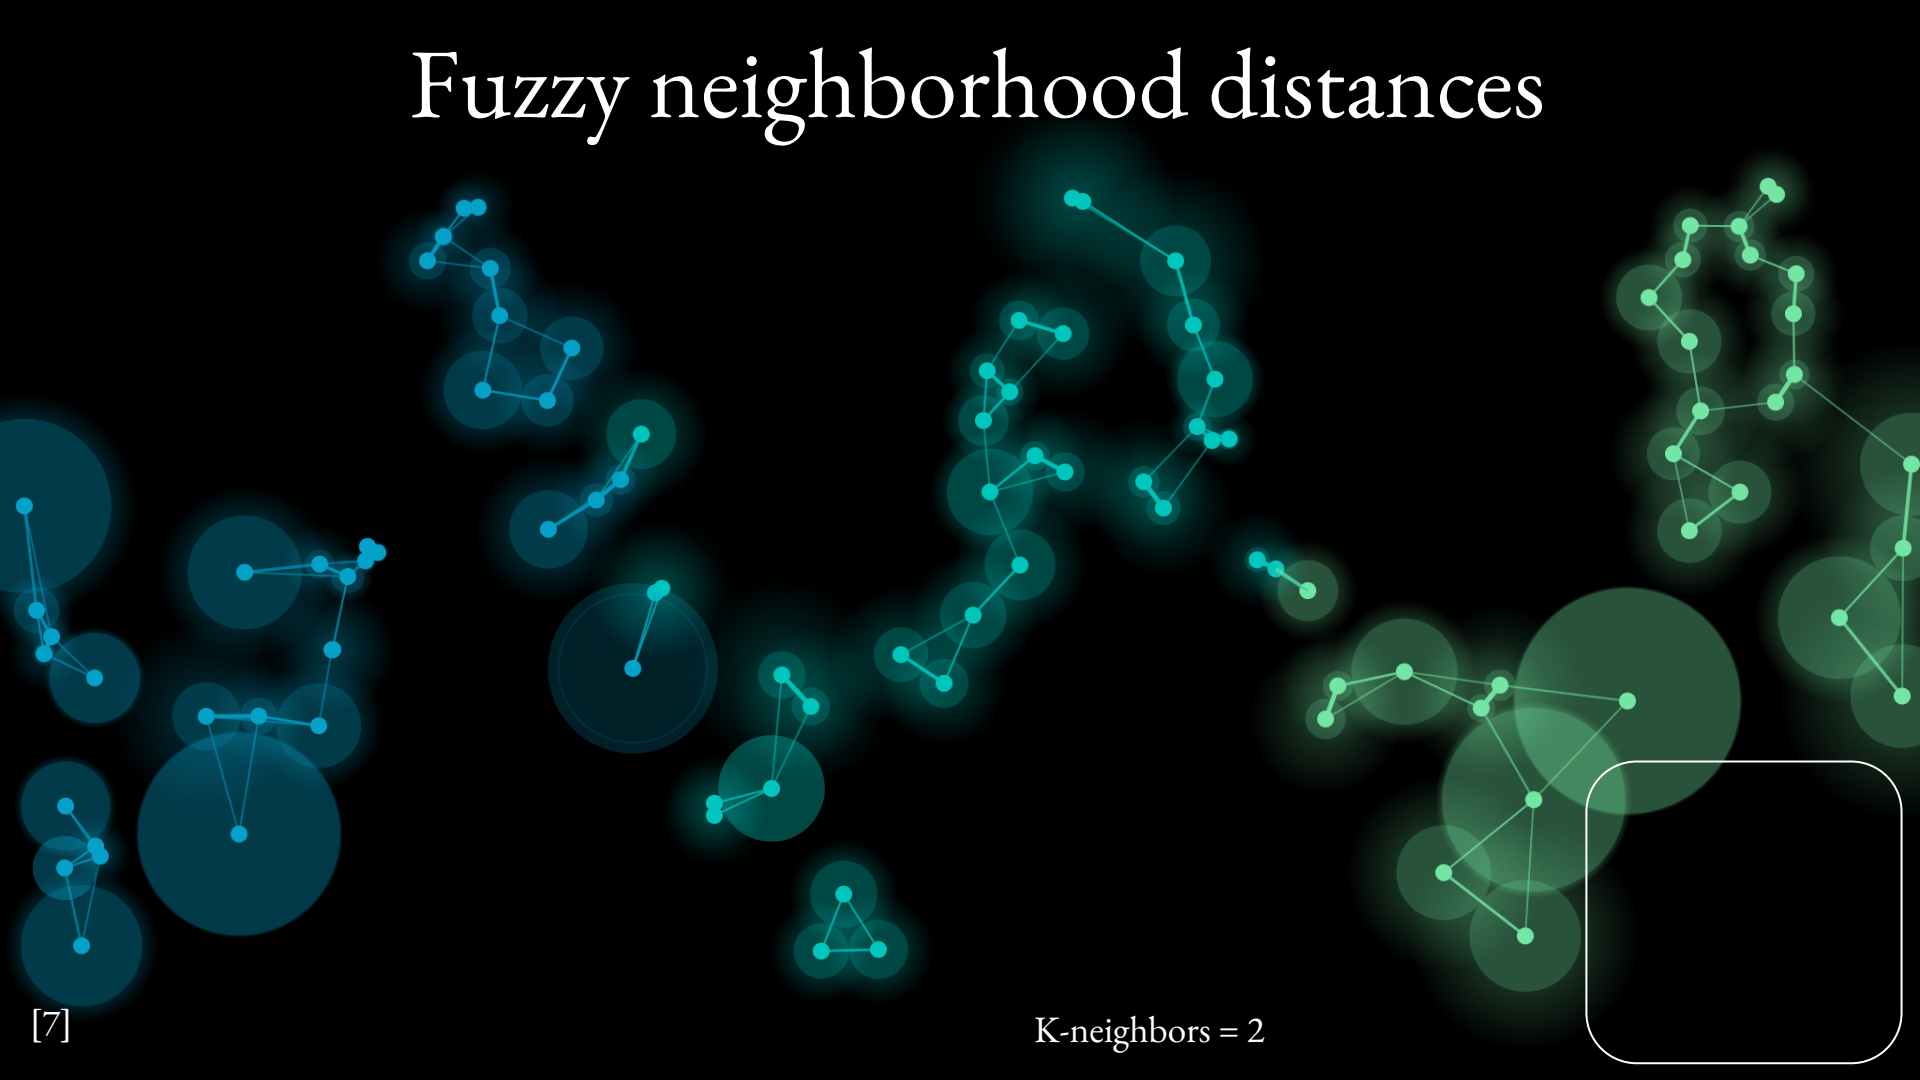

[7]

K-neighbors = 2

Fuzzy neighborhood distances

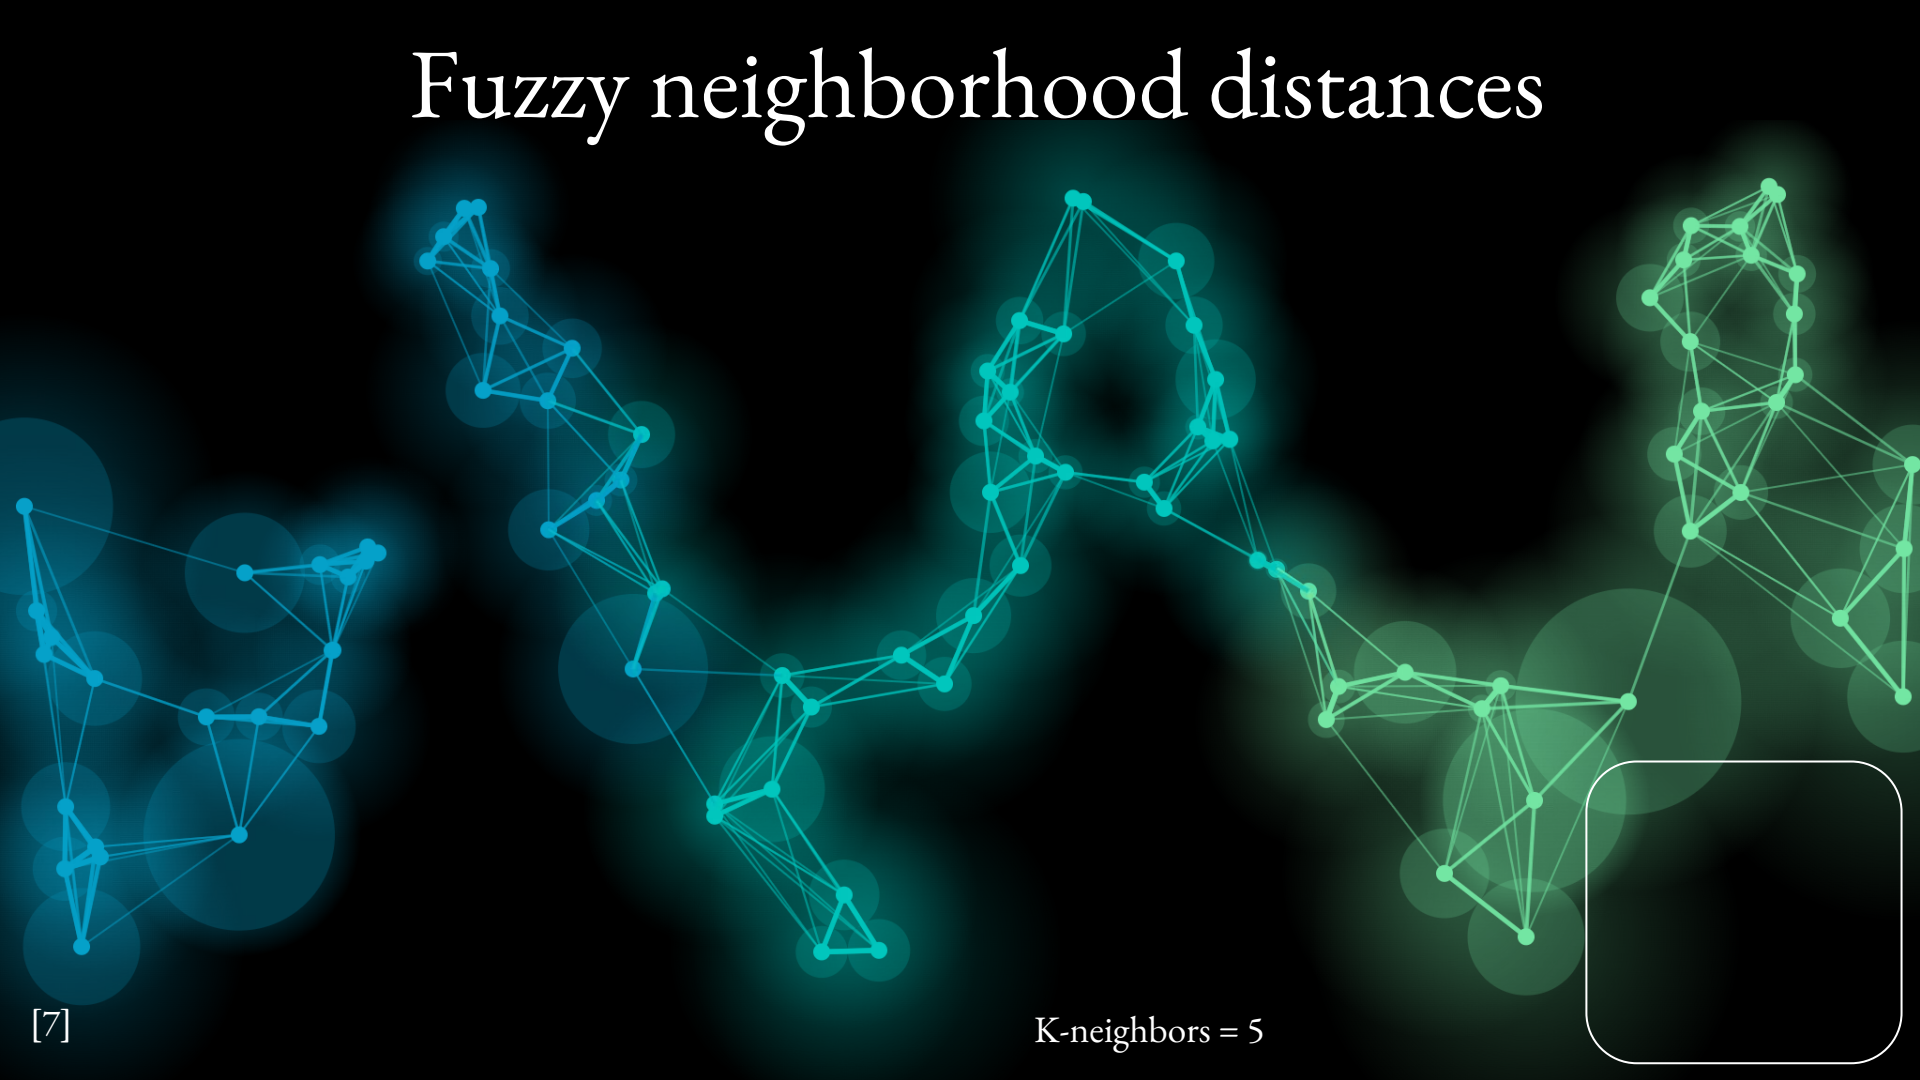

[7]

K-neighbors = 5

Fuzzy neighborhood distances

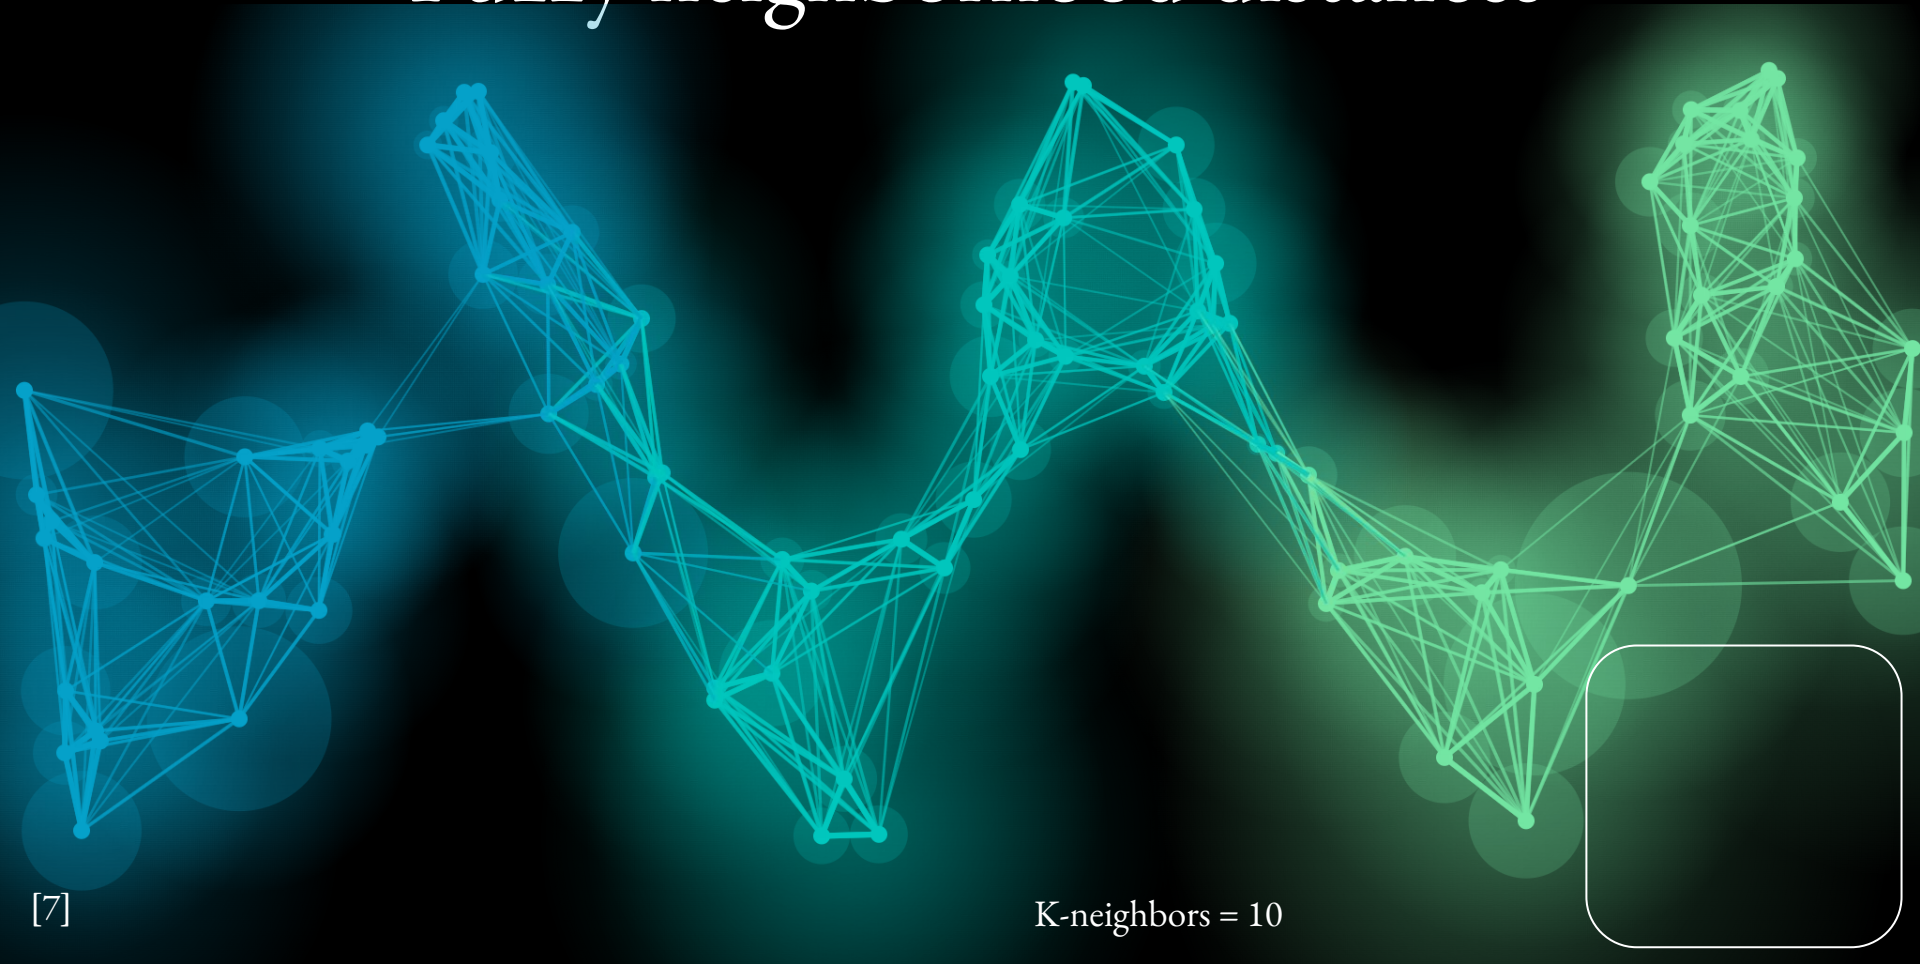

[7]

K-neighbors = 10

Initialize and optimize low-dimensional representation

Suppose the weight (probability of an edge $e \in E$ *is* $w_h(e)$ in the original (high dimensional) embedding and $w_l(e)$ in the new (low dimensional) embedding.

We want to minimize the cross entropy,

$$\sum_{e \in E} w_h(e) \log\left(\frac{w_h(e)}{w_l(e)}\right) + (1 - w_h(e)) \log\left(\frac{1 - w_h(e)}{1 - w_l(e)}\right)$$

Attractive Force

Repulsive Force

3. Practical Considerations

Practical Considerations

- Parameter choice can significantly impact runtime.
- Interpretation is tricky: orientation doesn't matter, sizes and distances of clusters don't mean much, and you may see patterns where none exist.
- The burden is on the user to pick parameters that fit your needs.
- It is tempting to use UMAP in places you should not.

4. Conclusion

To sum up:

- UMAP rocks.
- It provides (somewhat hefty) mathematical justification to get theoretically guaranteed results.
- It's a plug-in-place replacement for t-SNE and uniformly superior for visualizing high-dimensional data

Thanks for listening!

REFERENCES

[1] L. McInnes, J. Healy, and J. Melville, *Umap: Uniform manifold approximation and projection for dimension reduction,* arXiv preprint arXiv:1802.03426, 2018.

[2] L. McInnes and J. Healy, *How umap works*, How UMAP Works - umap 0.5 documentation. [Online]. Available: https://umap-learn.readthedocs.io/en/latest/how_umap_works.html

[3] D. Angelov, *Top2vec: Distributed representations of topics*, 2020.

[4] Wikipedia contributors, *Guernica (picasso) Wikipedia, the free encyclopedia*, 2021, [Online; accessed 17-November-2021]. [Online]. Available: https://en.wikipedia.org/w/index. php?title=Guernica_(Picasso)&oldid=1054068902

[5] D. I. Spivak, *Metric realization of fuzzy simplicial sets*, Self published notes, 2012.

[6] M. Noichl, *Examples for umap reduction using 3d models of prehistoric animals*, GitHub repository, 2019. [Online]. Available: https://github.com/MNoichl/UMAP-examples-mammoth\-

[7] A. Coenen and A. Pearce, *Understanding umap* [Online]. Available: https://pair-code.github.io/understanding-umap/

QUESTIONS

Please reach out to me at nicholasalines@gmail.com or drop by the github project!